

A Dormant Subspace, Not a Missing One: A Decodable but Causally Inert Answer-Site Representation of Arithmetic in Llama-3.1-8B

Anonymous ACL submission

Abstract

Large language models often fail at arithmetic composition even when they solve every part in isolation. We study a token-matched instance in Llama-3.1-8B: wrapping a multiplicative operand in an additive identity, $(0 + B) * C$, collapses multiplication accuracy (instruct 0.27, base 0.51) although the model evaluates the wrapped sub-expression perfectly (1.00) and multiplies inside a bracket (0.91). The failure is operation-specific (the additive analogue composes at 0.99), depth-sensitive, magnitude-controlled, and recovers with a few in-context examples. We then ask what the internal representation looks like. At the answer position the operand B and the near-product $B \cdot C$ are linearly decodable ($R^2 \approx 0.96$), which invites a “the value is there but unused” story. We show this reading is an interpretability illusion. The decode is operand-dominated: it has no margin over a linear- (B, C) baseline where there is no headroom, and only a small margin (gap ≈ 0.06 , CI excludes zero) where there is. More decisively, the decodable direction is causally inert: steering along it does not change the output, and a full-residual swap at the answer position does not flip it either, while patching the operand positions does move it (recovery 0.50), a dead site, not a dead probe. Decodability at the answer site is not evidence of computation there. The collapse replicates across model families in matched format, with a prompt-format confound we make explicit.

1 Introduction

Language models solve many arithmetic problems yet fail on superficially trivial variants. A recurring explanation is that some intermediate value is computed but not propagated, tested by probing the residual stream: if the value decodes, the model “knows” it. We show this inference is unsafe. A decodable value at the answer position can be an operand-reconstructible, causally inert shadow of the computation, not the computation itself.




Figure 1: The answer-site subspace is *dormant*, not missing. At the “=” position the operand and near-product are linearly decodable, but that direction is causally inert (steering and a full-residual swap change nothing) and operand-explained; the answer is instead re-derived diffusely from the operand positions. Decodability at the answer site is not evidence of computation there.

Our vehicle is a single controlled failure. For $B, C \in [20, 49]$, wrapping the first operand in an additive identity, $(0 + B) * C =$, collapses accuracy to 0.27 (instruct) / 0.51 (base) although the model returns B for $(0 + B) =$ at 1.00 and multiplies inside a bracket, $(B * C) =$, at 0.91; the gap is large and tight ($\Delta(C4 - C1) = 0.64$, 95% CI [0.60, 0.69], instruct). This is a clean hook, not our contribution: we use it to ask whether a decodable representation at the answer site licenses a claim about computation. Both B and the product $B \cdot C$ are linearly decodable there ($R^2 \approx 0.96$), inviting the reading that the product “is represented but not used.” Figure 1 summarizes why that reading fails.

Contributions.

- A token-matched, control-rich instance of a compositional arithmetic blind spot in Llama-3.1-8B: operation-specific, depth-sensitive, magnitude-controlled, and few-shot-recoverable, each with a confidence interval (§4).
- Evidence that the decodable answer-site subspace is *operand-dominated*: it has no margin over a linear- (B, C) baseline where the band leaves no headroom, and only a small margin (gap ≈ 0.06 , CI excludes zero) where it does, robust to a C1-independent probe-layer choice (§5).

- A causal test showing the decodable direction is *inert*: probe-direction steering and a full-residual swap at the answer position both leave the output unchanged, anchored by a positive control (patching the operand positions moves the answer, recovery 0.50), so the null is a dead site, not a dead probe. Ablating the decode direction preserves $\approx 85\%$ of the swap’s (sub-flip) answer-effect (§5).
- A cross-family replication in matched bare format, together with an explicit prompt-format confound that we argue precludes reading the cross-model results as a clean family dissociation (§6).

We read the answer-site subspace as *dormant* in the sense of Makelov et al. (2024): present to a probe, invisible to the computation.

2 Related Work

Surface form and mathematical brittleness.

Model math accuracy is sensitive to surface structure that leaves the answer invariant. Williamson et al. (2025), at the same venue, characterize a behavioral blind spot in natural-language word problems where failure tracks dependency-locality complexity and recovers under structure-lowering rephrasing, and conjecture a coupling between surface form and internal representation. Our instance is orthogonal on both axes and *tests* that conjecture: a token-matched *symbolic* expression whose additive-identity wrapper carries no dependency-locality cost, for which we show the obvious representational reading (a computed-but-unused product at the answer site) does not hold; the internal signature is a dormant, operand-explained subspace, not a missing representation. Reusch and Lehner (2022) show operator-tree structure is decodable from embeddings; we show that answer-site decodability of the arithmetic *value* can be an operand-reconstructible artifact that is causally inert, so decodability of arithmetic structure does not certify its use.

Decodability versus causal use. A linear probe measures correlation, not use (Belinkov, 2022), and can succeed on random labels without a control (Hewitt and Liang, 2019). The sharpest form is the subspace-patching illusion of Makelov et al. (2024): a direction can carry a decodable signal yet be causally disconnected from the output. Our result is an arithmetic instance of this illusion, es-

tablished with a baseline and a steering test rather than assumed. Matsumoto et al. (2022) find decoded intermediate values in neural math solvers to be causally usable; we find a decodable answer-site subspace that is causally inert. The results are compatible: whether a decodable value is used is a per-site empirical question.

Mechanisms of arithmetic. Stolfo et al. (2023) trace arithmetic from operand tokens to the final position through late modules, and Nikankin et al. (2025) argue models use a diffuse “bag of heuristics.” Our localization agrees, and adds that the decodable answer-site value, easily mistaken for the computed result, is a byproduct of that operand information. For method we follow patching best practice (Zhang and Nanda, 2024; Heimersheim and Nanda, 2024): symmetric counterfactuals, a full-answer log-probability metric, and a positive control per null.

3 Setup and Research Questions

Model and stimuli. We study Llama-3.1-8B (Grattafiori et al., 2024), base and instruction-tuned, in bare continuation format, by greedy decoding and integer parsing. Stimuli are $N = 400$ operand pairs (B, C) drawn once with a fixed seed from $[20, 49]$ and held constant across every condition, token-length matched, with spacing controlled. The central contrast is $C1 = (0 + B) * C =$ (the composed failure) against $C4 = (B * C) =$ (multiply inside a bracket, succeeds) and $C6 = (0 + B) =$ (evaluate the wrapper, succeeds).

Metrics. Accuracies carry Wilson intervals; paired contrasts carry paired bootstrap intervals (10,000 resamples). Decodability is five-fold cross-validated ridge R^2 (item-bootstrap interval; selectivity-gap CIs use 300 resamples given the per-draw cross-validation cost) at a named token position. Selectivity compares the decode against a linear- (B, C) baseline (a probe given only the raw operands and Gaussian noise), so a high R^2 that is merely linear in the operands does not count as representing the product. Causal tests use symmetric operand counterfactuals and a full-answer log-probability metric $\Delta = \log p(P') - \log p(\text{true})$; because Llama chunks the leading digits, a first-token score is non-discriminative between two products.

Research questions. **RQ1:** is the composition failure a genuine, controlled effect or an artifact of difficulty, magnitude, or elicitation (§4)? **RQ2:**

	surface	base	instruct
C0	$B * C =$	0.838	0.848
C4	$(B * C) =$	0.890	0.908
C6	$(0 + B) =$	1.000	1.000
C1	$(0 + B) * C =$	0.508	0.265
A1	$(0 + B) + C =$	n/a	0.995
D1	$((0 + B)) * C =$	n/a	0.038

Table 1: The composition blind spot (exact-match accuracy, $N=400$). The parts succeed (C4, C6) but the composition collapses (C1); $\Delta(C4-C1) = 0.64$, CI [.60, .69] (instruct). A1/D1 are instruct-only. Full battery with Wilson intervals in Appendix A.

is the near-product decodable at the answer site a *represented* quantity or reconstructible from the operands alone (§5)? **RQ3:** is the decodable direction causally used to produce the answer (§5)?

4 The Blind Spot

The model evaluates every ingredient of $(0 + B) * C$ yet fails to compose them (Table 1). We establish that this is a real effect, not an artifact, with four controls, and we state each at the precision the data support.

It is operation-specific, not format brittleness. Replacing the outer multiplication with addition, $(0 + B) + C =$, restores accuracy to 0.995 (instruct) against 0.265 under *identical* format; $\Delta(A1 - C1) = 0.730$, 95% CI [0.69, 0.77]. It is also spacing-invariant: with all spaces removed the inside-bracket multiply still survives (C8 = 0.785 instruct, 0.765 base) while the composition still collapses (C7 = 0.19, 0.018), the same dissociation. So it is not generic spacing or parenthesis sensitivity; specifically the multiplicative composition breaks. (A1/A2 are instruct-only; we do not extrapolate to base.)

It is controlled, not an artifact. Three further controls hold. *Depth* is a distinct stressor: a redundant nesting $((0 + B)) * C =$ drops to 0.038, but nesting degrades *addition* too (instruct nested addition falls to 0.25 from 0.99), so we treat depth as a general parsing stressor, not more composition evidence. *Magnitude* is ruled out: within every matched- $|B \cdot C|$ bin C1 stays well below C4 (0.53 vs 0.96 smallest, 0.00 vs 0.70 largest). And the collapse is an *elicitation* fragility: two to four fixed-seed in-context examples recover C1 to the C4 ceiling (instruct 0.265 $\rightarrow$ 0.885 [.85, .91] $\rightarrow$ 0.915 [.88, .94]; base 0.508 $\rightarrow$ 0.838 $\rightarrow$ 0.882), so it is not an inability to multiply.

Two overstatements we avoid. The effect is neither universal (base $(B + 0) * C$ barely moves, drop 0.015, vs 0.33 for $(0 + B)$; Appendix B) nor uniformly worse under tuning (on plain parentheses instruct is *better*, 0.755 vs 0.495); and C0, not 1.0, is the reference ceiling ($B * C =$ is 0.848 instruct).

5 A Dormant Subspace

The seductive story is that the product is computed and represented at the answer position but not read out. We test it in three steps and reject it.

The product decodes (RQ2, correlational). A cross-validated probe recovers $B \cdot C$ from the “=” residual at $R^2 = 0.965$ (instruct), 0.967 (base) and the operand B at 0.96–0.97, with B decodable at “)” at *every* layer including the last (instruct 0.98, layer 31); taken alone this looks like a represented product.

But the decode is operand-dominated (RQ2, selectivity). Decodability of BC need not mean the product is represented, since $B \cdot C$ is itself largely linear in B, C over a narrow band. We compare against a linear- (B, C) baseline (Figure 2a). At [20, 49] the baseline reaches $R^2 = 0.968$, equal to the decode, so the gap is -0.0002 : no headroom, nothing beyond the operands. At [2, 99], where BC is nonlinear in the operands, a small gap appears: $+0.068$, 95% CI [0.051, 0.086] (base) and $+0.060$, [0.043, 0.079] (instruct), both excluding zero, and still positive at a layer chosen by a C1-independent criterion (the C4 decode peak) rather than the C1 argmax. We flag that [2, 99] is a harder regime where the parts fall below our 0.80 floor (C4 = 0.71 base, 0.73 instruct), so we read it as a stress test, not a comfortably in-scope measurement, and rest the headline on the band-independent causal null below. The decode is operand-dominated (≥ 0.94 of it is reproduced by the linear- (B, C) baseline, ≈ 1.0 at [20, 49]) with at most a small operand-nonlinear component, consistent with a weak product signal (behavior-tracking: C4 gap > C1 gap, difference CI excludes zero).

And the decodable direction is causally inert (RQ3). We steer the “=” residual along the fitted direction to write a counterfactual product P' , on held-out items. On the first-token metric the inject is null at every layer and matches norm-matched random and shuffled controls (Appendix D), but that metric is non-discriminative be-




Figure 2: The decodable answer-site subspace is dormant (bars are 95% CIs). **(a)** Selectivity gap (decode R^2 for $B \cdot C$ at “=” minus a linear- (B, C) baseline): at [20, 49] no headroom (gap ≈ 0 , operand-explained); at [2, 99] a small gap whose CI excludes zero. **(b)** Causal share of a full-residual swap at the C4 “=” (base): ablating the decode direction (perp) preserves most of the effect, the decode direction alone (w-only) reproduces little, CIs non-overlapping; the swap flips no answers (emit- $P' = 0$). **(c)** The dissociation as a fraction of the answer moved (two metrics): patching moves the answer at the operand positions (recovery 0.50) but the answer-site swap flips nothing, a dead *site*, not a dead probe.

tween products. On the discriminative full-product metric, injecting at the C1 “=” is flat at every dose $k \in \{1, 2, 4, 8\}$ (CIs bracket zero) while the same inject at the succeeding C4 “=” rises monotonically ($+0.08 \rightarrow +0.46$, every CI excluding zero, base) yet never flips: the axis has dynamic range where the product is used and none at the failing site, a dead direction rather than a dead instrument (emit- P' stays zero throughout, a saturating criterion, so we read the log-probability). A full-residual swap at the C4 “=”, where the model succeeds, likewise fails to transplant the donor product (emit- $P' = 0$; log-probability $+0.59$ at L19 where we decompose it, peak $+0.66$ at L23, still sub-flip), so the product is not linearly resident at “=” even where computed correctly, a fortiori at C1. The one intervention that moves the answer is at the operand positions (recovery 0.50; a clean-answer recovery metric, a different readout and operation).

Certifying the illusion. We decompose the C4 swap effect into the decode direction and its orthogonal complement (Figure 2b). Ablating the decode direction still leaves $+0.50$ (base), while the decode direction alone contributes only $+0.08$, with non-overlapping CIs: $\approx 85\%$ of the swap’s (sub-flip) answer-effect survives removing the decodable direction (whose magnitude is only $\approx 20\%$ along that direction). (The instruct swap effect is smaller, $+0.12$, its decode-direction component not significantly positive; we report the decomposition for base.) We use *dormant* strongly: unlike a patch-potent-but-unused component, this direction moves essentially nothing even when patched, present to a probe but invisible to computation (Makelov et al., 2024). Consistent with Stolfo et al. (2023) and Nikankin et al. (2025), the operands carry the signal but no sparse head set reconstructs it (best head < 0.02); the value at “=” decodes and correlates

yet cannot be steered.

6 Generality and Limitations

Replication in matched format, and a format confound. In matched bare format the collapse is not specific to Llama-3.1-8B. Among models that clear a parts floor (C4, $C6 \geq 0.80$), the composed form collapses for Gemma-2-9b-it (C1 = 0.043, C4 = 0.99) and Llama-3.2-3B (C1 = 0.36), alongside base and instruct Llama-3.1-8B; three others fail the parts in bare format and are out of scope, not counterexamples (Appendix E). Chat formatting changes this, not cleanly: Gemma-2-9b-it *composes* in its native chat template (C1 = 0.95), but there Llama-3.1-8B falls below the parts floor (C4 = 0.70) and leaves scope, so the composing and failing models are never compared in the same format. We read this as replication in matched bare format plus a prompt-format confound, *not* as evidence that some families are immune; per-model native-format evaluation is the right axis for a clean cross-family claim, which we do not make.

Limitations. The causal localization is distributed (operand positions matter but no sparse circuit is recovered), so we characterize *where*, not a minimal mechanism; the operation and depth controls are instruct-only; and the zero-excluding gap is small and only in the harder band, with the dormancy resting on the causal-share certification. None of these affects the central result.

Conclusion. A parenthesized additive identity selectively breaks multiplicative composition in Llama-3.1-8B, and the “represented but unused” reading of the decodable answer site is an illusion: it is operand-dominated and causally inert, so a high decodability R^2 needs a baseline and a causal control before it licenses a mechanistic claim.

References

- Yonatan Belinkov. 2022. [Probing classifiers: Promises, shortcomings, and advances](#). *Computational Linguistics*, 48(1):207–219.
- Aaron Grattafiori et al. 2024. [The Llama 3 herd of models](#). *Preprint*, arXiv:2407.21783.
- Stefan Heimersheim and Neel Nanda. 2024. [How to use and interpret activation patching](#). *Preprint*, arXiv:2404.15255.
- John Hewitt and Percy Liang. 2019. [Designing and interpreting probes with control tasks](#). In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2733–2743, Hong Kong, China. Association for Computational Linguistics.
- Aleksandar Makelov, Georg Lange, and Neel Nanda. 2024. [Is this the subspace you are looking for? an interpretability illusion for subspace activation patching](#). In *The Twelfth International Conference on Learning Representations (ICLR)*.
- Yuta Matsumoto, Benjamin Heinzerling, Masashi Yoshikawa, and Kentaro Inui. 2022. [Tracing and manipulating intermediate values in neural math problem solvers](#). In *Proceedings of the 1st Workshop on Mathematical Natural Language Processing (MathNLP)*, pages 1–6, Abu Dhabi, United Arab Emirates (Hybrid). Association for Computational Linguistics.
- Yaniv Nikankin, Anja Reusch, Aaron Mueller, and Yonatan Belinkov. 2025. [Arithmetic without algorithms: Language models solve math with a bag of heuristics](#). In *The Thirteenth International Conference on Learning Representations (ICLR)*.
- Anja Reusch and Wolfgang Lehner. 2022. [Extracting operator trees from model embeddings](#). In *Proceedings of the 1st Workshop on Mathematical Natural Language Processing (MathNLP)*, pages 40–50, Abu Dhabi, United Arab Emirates (Hybrid). Association for Computational Linguistics.
- Alessandro Stolfo, Yonatan Belinkov, and Mrinmaya Sachan. 2023. [A mechanistic interpretation of arithmetic reasoning in language models using causal mediation analysis](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 7035–7052, Singapore. Association for Computational Linguistics.
- Dane A Williamson, Yangfeng Ji, and Matthew B. Dwyer. 2025. [Syntactic blind spots: How misalignment leads to LLMs’ mathematical errors](#). In *Proceedings of The 3rd Workshop on Mathematical Natural Language Processing (MathNLP 2025)*, pages 1–14, Suzhou, China. Association for Computational Linguistics.

Fred Zhang and Neel Nanda. 2024. [Towards best practices of activation patching in language models: Metrics and methods](#). In *The Twelfth International Conference on Learning Representations (ICLR)*.

A Full Battery and Confidence Intervals

Table 2 gives every condition with Wilson 95% intervals ($N = 400$). Base A1/A2/D1 are not in our committed data and are omitted rather than estimated.

surface	base	instruct
C0 $B * C =$	.838 [.80,.87]	.848 [.81,.88]
C1 $(\emptyset + B) * C =$	.508 [.46,.56]	.265 [.22,.31]
C2 $(\emptyset + B * C) =$	n/a	.678 [.63,.72]
C3 $(B) * C =$	n/a	.755 [.71,.80]
C4 $(B * C) =$	.890 [.86,.92]	.908 [.88,.93]
C5 $\emptyset + B * C =$	n/a	.542 [.49,.59]
C6 $(\emptyset + B) =$	1.00 [.99,1.0]	1.00 [.99,1.0]
C7 $(\emptyset + B) * C =$	.018 [.01,.04]	.190 [.16,.23]
C8 $(B * C) =$	.765 [.72,.80]	.785 [.74,.82]
A1 $(\emptyset + B) + C =$	n/a	.995 [.98,1.0]
A2 $\emptyset + (B + C) =$	n/a	.928 [.90,.95]
D1 $((\emptyset + B)) * C =$	n/a	.038 [.02,.06]

Table 2: Full battery, exact-match accuracy with Wilson 95% CIs.

Magnitude control. Within matched $|B \cdot C|$ bins (instruct), C1 vs C4: [440, 813] .53/.96 ($n=79$); [813, 1185] .28/.94 (126); [1185, 1558] .16/.89 (104); [1558, 1930] .19/.89 (64); [1930, 2303] .00/.70 (27). Every paired difference has a positive bootstrap CI.

B Wrapper Blind-Spot Map




Figure 3: Multiplication accuracy by wrapper. Position-dependence is a base-model effect (left-placed $(\emptyset + B)$ collapses, right-placed $(B + \emptyset)$ barely moves; instruct degrades on both placements), and instruction tuning is not uniformly worse (plain parentheses).




Figure 4: Operand B decodes at every layer in the failing zero-shot regime (solid). The few-shot curves (dashed) drop mid-network; shown for completeness only, and we draw no mechanistic conclusion from the few-shot decodability.

C By-Layer Decodability

400

D Steering Nulls

401




Figure 5: Probe-direction steering at “=” is inert at every layer for inject, random, and shuffled targets (mean Δ ground-truth-token logit, 95% CIs through zero). This first-token metric is a null; the discriminative full-product re-score is Figure 2b.

E Cross-Model Battery

402

Matched bare format, in scope if C4 and C6 $\geq$ 0.80: Llama-3.1-8B base C1 .508 / instruct .265; Gemma-2-9b-it .043 (C4 .99); Llama-3.2-3B .358 (C4 .88). Out of scope (parts fail): Qwen2.5-7B (C4 .00), Mistral-7B (C4 .62), Llama-3.2-1B (C4 .31). Chat format: Gemma-2-9b-it composes (C1 .95, C4 .99) but Llama-3.1-8B chat falls out of scope (C4 .70); the composing and failing models are never in the same format.

403

404

405

406

407

408

409

410

411